University of New South Wales

COMP9417 - Group Project - xRFM

Group: Twinkle Twinkle

John Zhang	z5591121
Shaohua Wang	z5489497
Zian Hu	z5514488
Yonje Ryu	z5338581

1. Introduction

Tabular data resists deep learning: targets jump abruptly along feature axes, scales are heterogeneous, and many columns are uninformative; gradient-boosted decision trees (GBDTs) handle these properties naturally [Grinsztajn et al., 2022]. Classical kernel methods don't displace them either: kernels fix feature similarity *before* training (so they cannot learn which features matter), and kernel ridge regression costs $O(n^3)$, intractable above ~70k samples.

xRFM [Beaglehole et al., 2025] tackles both. An inner Recursive Feature Machine [Radhakrishnan et al., 2024] learns the Average Gradient Outer Product (AGOP), a matrix of per-input sensitivities, and reshapes the kernel through it. An outer balanced binary tree splits along the top AGOP eigenvector of a local RFM, bringing training to $O(n \log n)$ and producing a per-leaf AGOP for subpopulation interpretability. We test xRFM on five UCI datasets (none in the TALENT benchmark or xRFM's meta-test suite) against XGBoost, Random Forest, and CatBoost. xRFM is competitive on smooth-target tasks (Appliances Energy, Crop Mapping), trails on rule-like and small-n problems (Seoul Bike, HCC), and is more hyperparameter-sensitive on imbalanced high-dimensional data (IDA2016).

2. Methodology

2.1 Average Gradient Outer Product (AGOP)

AGOP is the matrix at the heart of xRFM's feature learning: across training points, in which input directions does the model's prediction change most? For a trained predictor $\hat{f} : \mathbb{R}^d \to \mathbb{R}$ and training set $\mathcal{S} = \{x^{(1)}, \dots, x^{(n)}\}$, AGOP is computed by taking the gradient of $\hat{f}$ at every training point, then averaging the outer products:

$$\mathrm{AGOP}(\hat{f}, \mathcal{S}) \;=\; \frac{1}{n} \sum_{i=1}^{n} \nabla \hat{f}(x^{(i)})\, \nabla \hat{f}(x^{(i)})^{\top} \;\in\; \mathbb{R}^{d \times d}. \tag{1}$$

For any unit vector v, the quadratic form $v^\top M v$ measures how strongly $\hat{f}$ responds when its input moves along direction v. Two readouts of M answer different questions:

- **Diagonal entries** M_{ii} give per-feature sensitivity. A large M_{ii} means the model's output changes sharply when feature i changes, holding the others fixed.
- **Top eigenvectors of** M give the directions in feature space that the model actually uses (the *active subspace* of Constantine et al., 2014). The top eigenvector v_1 is a signed linear combination of features (positive and negative loadings push the prediction in opposite directions), capturing *joint* structure that the diagonal alone misses.

How AGOP compares to other importance methods. *PCA* finds high-variance directions of x but ignores y, so its directions need not be predictive. *Mutual information* $I(X_i; Y)$ is model-free and per-

feature interpretable but cannot capture joint structure. *Permutation importance* shuffles a feature and measures the loss change; intuitive but degrades under correlated features. AGOP is supervised (depends on $\hat{f}$), joint (off-diagonal entries encode interactions), and computed in a single forward/backward pass; the cost is that it inherits the model's biases: if $\hat{f}$ overfits noise, AGOP assigns it apparent importance.

2.2 The xRFM Algorithm

xRFM nests an iterative kernel-ridge RFM inside the leaves of a balanced binary tree whose internal splits use an AGOP-based supervised criterion.

Tree construction. Given a node with more than $L_{\max}$ points, xRFM trains a small RFM on a subsample and takes the top eigenvector v_1 of the local AGOP as the split direction: every point is projected onto v_1 and split at the **median**, recursing until every leaf has at most $L_{\max}$ points. Median splits keep the tree balanced (depth $O(\log n)$), which is the design choice distinguishing xRFM from prior AGOP-based trees and the source of its $O(n \log n)$ training cost.

Leaf training. At each leaf, an RFM is fitted using the generalised kernel $K_{p,q}(x, x') = \exp(-\|x - x'\|_p^q / L^q)$ (Schoenberg, 1942) for $0 < q \leq p \leq 2$. Two cases matter for our experiments: $q = 2$ recovers the rotation-invariant Laplace kernel (the library's `l2` mode), and $q = p$ gives an axis-aligned product kernel (`l1`). Starting with $M_0 = I$, the RFM iteration alternates kernel-ridge fitting on inputs reshaped by $M_t^{1/2}$ with AGOP recomputation and normalisation, for up to $\tau \leq 10$ rounds; we retain the iterate with the lowest *validation* error.

Inference. A test point $x^\star$ is routed through the tree (hard or soft median thresholds) to its leaf ℓ, then predicted as $\hat{y} = K(x^\star \odot M_\ell^{1/2}, X_{M_\ell})\alpha_\ell$. A single tree, with one predictor per test point and no ensembling.

Computational complexity. Training is $O(n \log n)$: leaves cap per-leaf cost at $O(L_{\max}^3)$ (constant in n), and median splits visit $O(n)$ points at each of $O(\log n)$ tree levels. Inference is $O(\log n + L_{\max})$ per sample. A naive kernel method costs $O(n^3)$ training and $O(n^2)$ memory, which becomes intractable above roughly 70,000 samples on a 40 GB GPU.

2.3 Experimental Design

Datasets. Five UCI datasets (Table 1), absent from both the TALENT benchmark and xRFM's meta-test suite. The selection spans stress points for kernel methods: smooth-target regression (Appliances, Crop Mapping), rule-like targets (Seoul Bike), small-n high-d/n with missingness (HCC), and severe class imbalance (IDA2016, ~1.6% positive). Appliances doubles as the interpretability testbed (ships with two ground-truth noise features, `rv1`/`rv2`).

#	Dataset (UCI ID)	n	d	Task	Features
1	Seoul Bike Sharing (560)	8,760	12	Regression	9 num + 3 cat
2	Appliances Energy (374)	19,735	27	Regression	All numeric
3	HCC Survival (423)	165	49	Binary	22 num + 27 cat
4	IDA2016 Scania APS (414)	60,000	170	Binary	All numeric
5	Crop Mapping, Winnipeg (525)	50,000	174	Multiclass (7)	All numeric

Splits. Stratified 60/20/20 train/val/test with `random_state=42` (stratified on target for classification, random for regression).

Preprocessing. Median imputation for numerics, mode imputation for categoricals. For kernel methods (xRFM): standard-scaled numerics + one-hot-encoded categoricals. For tree ensembles: numerics plus ordinal-encoded categoricals (no standardisation, preserving split structure). CatBoost receives native categorical indices.

Hyperparameter tuning. Optuna TPE sampler optimises the validation-split primary metric (RMSE for regression, AUC-ROC for binary, accuracy for multiclass) with model-specific trial budgets (xRFM 25, XGBoost 40, CatBoost 30, Random Forest 20). xRFM gets fewer trials because each kernel-ridge solve is the wall-clock bottleneck. The configuration with the best validation objective is retrained on the train split and evaluated once on held-out test. Full search spaces, fixed knobs, and per-dataset selected HPs are in Appendix B. Note that `max_leaf_size = min(60000, n_train)` forces xRFM into a single-leaf-RFM regime on all five main-comparison datasets; the tree component is engaged only in the §3.2 interpretability run (2 leaves on Appliances) and at $n = 36k$ in §3.3.

Metrics. Regression: RMSE (primary) and R^2. Classification: Accuracy and AUC-ROC (macro-OvR for multiclass). Plus wall-clock training time and per-sample inference time.

3. Results

3.1 Main comparison

Table 2. Final test-set results across all five datasets and four models. Regression: lower RMSE is better. Classification: higher Acc/AUC is better. Train is wall-clock training time (s); Inf is per-sample inference time (µs). Best per dataset in bold; ties broken by AUC for binary, Accuracy for multiclass.

Dataset	Metric	xRFM	XGBoost	CatBoost	RF
Appliances Energy	RMSE/R²	**65.29/0.574**	65.75/0.568	66.96/0.552	72.61/0.473
	Train/Inf	11.5/26	41.7/95	11.1/**3.5**	**1.7**/41
Seoul Bike	RMSE/R²	238.44/0.864	232.32/0.870	**225.66/0.878**	245.64/0.855
	Train/Inf	2.1/16	5.4/37	5.2/**1.9**	**0.5**/57

Dataset	Metric	xRFM	XGBoost	CatBoost	RF
HCC Survival	Acc/AUC	0.758/0.800	0.818/0.835	0.788/0.819	**0.849/0.850**
	Train/Inf	0.1/50	**0.0/19**	0.6/47	0.9/3667
IDA2016	Acc/AUC	0.990/0.990	0.993/**0.993**	**0.994**/0.992	0.992/0.989
	Train/Inf	16.9/5.9	52.1/30	**3.1/0.5**	5.8/15
Crop Mapping	Acc/AUC	**0.995/1.000**	0.991/1.000	n/a	0.989/1.000
	Train/Inf	12.8/**2.6**	40.0/20	n/a	**12.2**/24

Seeds fixed at 42. TabPFN-v2 omitted throughout (see §4 for rationale). CatBoost on Crop Mapping timed out on our 2-hour Modal budget at 7-class × 30k × 174-feature scale; other three models converge.

Tuning chose `kernel=l1` (axis-aligned) for the two lowest-d datasets (Appliances, Seoul Bike) and `kernel=l2` (rotation-invariant) for the three higher-d datasets (HCC, IDA2016, Crop Mapping); per-dataset HPs are in Appendix B.

xRFM is competitive on smooth-target tasks. On Appliances Energy it obtains the best RMSE in this single split, but the margin over XGBoost is small (65.29 vs 65.75) and should not be over-interpreted without multi-seed error bars. On Crop Mapping, xRFM is best among the completed models, but CatBoost timed out, so this is not evidence that xRFM dominates all GBDT baselines. **GBDTs win on rule-like Seoul Bike** (CatBoost leads, all four within 9%; zero rentals when the system is off, and season × hour interactions favour splits) **and small-n HCC** (RF leads xRFM by 9 accuracy points; with 33 test samples, differences under ±6 points aren't significant at 5%).

3.2 Interpretability on Appliances Energy

We compare AGOP from the tuned xRFM (`kernel=l1, diag=True`) against PCA loadings (unsupervised), mutual information (model-free supervised), and signed permutation importance on an XGBoost proxy (black-box supervised). `diag=True` restricts AGOP to per-feature sensitivities; signed-eigenvector analysis (Fig. 7C) needs `diag=False`, deferred.

The random-feature test. Out of $d = 27$ features, MI ranks `rv1`, `rv2` at 26/27 and 27/27 and signed permutation at 27/27 and 26/27, so both cleanly identify the noise. AGOP places them at 21/27 and 20/27, correct in direction but less decisive. PCA fails outright at **2/27 and 1/27** (wrongly at the top), because random uniforms are statistically independent of the correlated sensor readings and so dominate input covariance. AGOP's noise rejection also depends on the kernel fit: re-running with default HPs (`bandwidth=10, kernel=l2, diag=False`) drops `rv1`, `rv2` to ranks 12–13, a nine-position degradation.

Per-leaf subpopulation structure. The tuned tree has two leaves; AGOP reveals different locally-relevant humidity sensors per leaf: leaf 0 weights RH_6 (kitchen) maximally (AGOP = 1.00), leaf 1 weights RH_8 (parents' room) maximally, mirroring the per-leaf specialisation Beaglehole et al. (2025, Fig. 7B) report for NYC Taxi. T6 (outdoor) and RH_5 (bathroom) rank high in both. With only 2

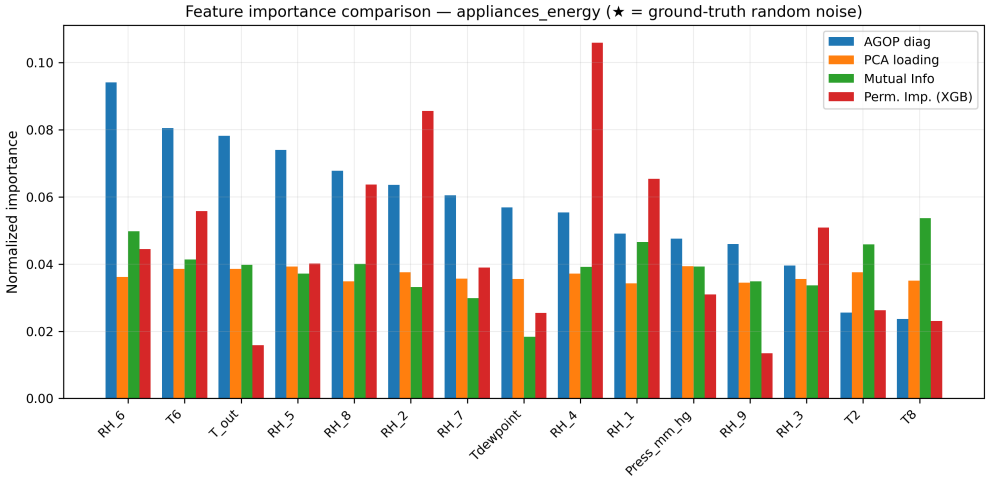

Figure 1: Feature importance comparison on Appliances Energy (top-15 features by AGOP). [*] marks the two injected random features.

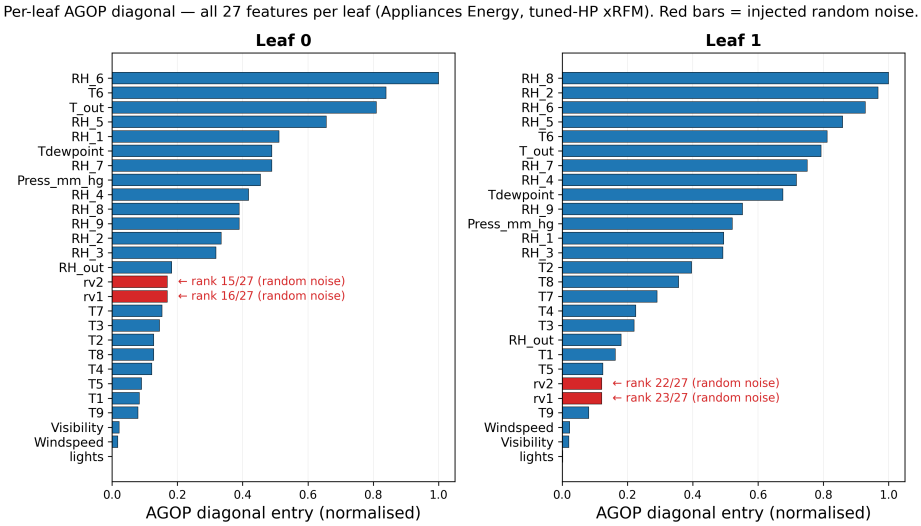

Figure 2: Per-leaf AGOP diagonal for the tuned-HP xRFM. The two leaves identify different humidity sensors as locally most important: RH_6 in leaf 0 and RH_8 in leaf 1.

leaves the effect is small-scale; deeper trees on larger datasets would amplify it. Per-leaf `rv1`/`rv2` ranks are 15–16/27 in leaf 0 and 22–23/27 in leaf 1, so noise demotion strengthens with deeper splits.

AGOP and signed permutation also agree on which real features matter: four of permutation's top-5 (RH_4, RH_2, RH_8, T4) appear in AGOP's top-10, both highlighting indoor humidity over outdoor temperature.

3.3 Scaling study on IDA2016

Figure 3 plots test AUC-ROC and training time against subsampled training size $n \in \{1k, 2.5k, 5k, 10k, 20k, 36k\}$ with the test set held fixed. Unlike the earlier default-HP sweep, this rerun freezes each model's Optuna-selected full-n hyperparameters from §3.1 and refits preprocessing on each subsampled training set only.

The AUC panel uses a broken y-axis so the XGBoost $n = 1k$ chance-level point remains visible while the 0.98+ comparison band is readable. This better isolates the effect of training-set size than the default-HP sweep, but it is still single-seed and should not be read as an asymptotic complexity test.

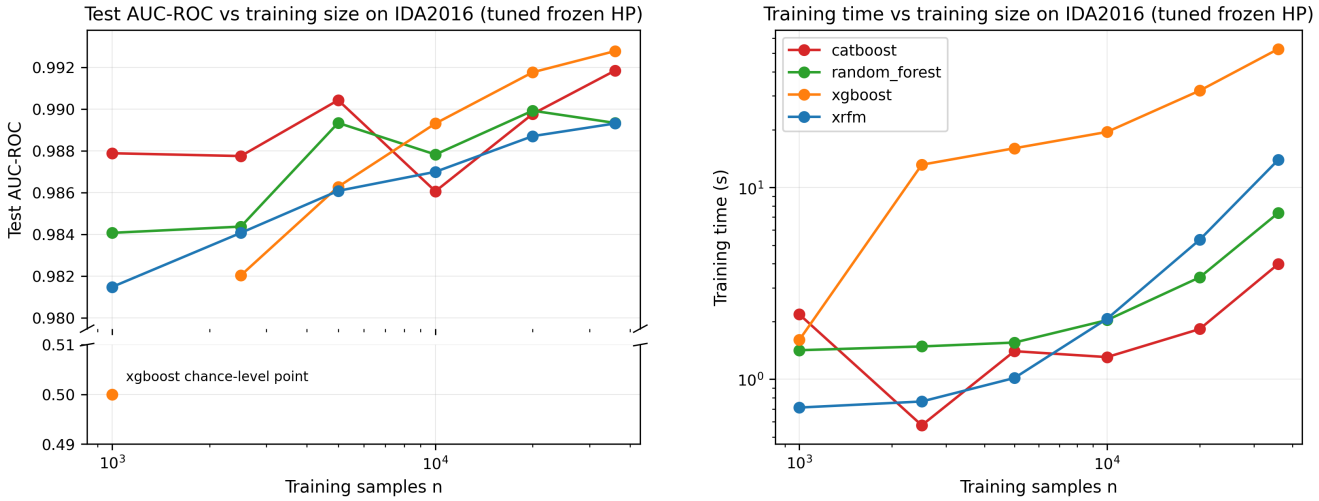

Figure 3: Test AUC-ROC with broken y-axis (left) and training time (right) vs. training size on IDA2016 with full-n tuned hyperparameters frozen.

1. **Tuned sample scaling.** xRFM improves AUC from 0.981 ($n = 1k$) to 0.989 ($n = 36k$), matching RF (0.989) but trailing CatBoost (0.992) and XGBoost (0.993). XGBoost's full-n tuned configuration fails at $n = 1k$ (AUC 0.500), showing frozen-HP scaling is configuration-specific.
2. **Tree engagement is shallow.** xRFM logs one leaf through $n = 20k$ and two at $n = 36k$ (the library rescales effective `max_leaf_size` to ≈26.6k); the sweep mostly measures leaf-RFM behaviour.
3. **Training time.** xRFM grows from 0.71 s ($n = 1k$) to 13.93 s ($n = 36k$); log-log slope 0.84. Faster than XGBoost (52.7 s) but slower than CatBoost (4.0 s) and RF (7.4 s) at full size; these are wall-clock measurements, not a complexity-claim test.

4. Discussion

Theory alignment. xRFM should be strongest on smoothly-varying targets and weaker on rule-like ones where threshold splits are the right inductive bias. The results are broadly consistent with this: xRFM is highly competitive on Appliances Energy and Crop Mapping (continuous sensors / 174 spectral bands), and loses narrowly on Seoul Bike (sharp on/off rules, season × hour interactions). The surprise is HCC Survival. Classical kernel theory predicts that kernels should do well in small-n regimes via RKHS regularisation; instead Random Forest beats xRFM by 9 accuracy points. Our hypothesis: with $n_{\text{train}} = 99$, learning a 108×108 AGOP matrix (the encoded dimension after one-hot expansion of the 27 categorical features) is statistically fragile, so noisy AGOP implies noisy reweighting.

xRFM is more hyperparameter-sensitive than the GBDTs. IDA2016 makes this clearest: tuned xRFM reaches AUC 0.990 (within 0.3 of CatBoost/XGBoost) but the earlier default-HP run gave only

0.978, a 1.2-point configuration effect. Three IDA2016 features match Grinsztajn et al. (2022)'s kernel-loss predictions:

- **Class imbalance (1.6% positive).** Kernel ridge minimises squared error symmetrically, with no minority up-weighting.
- **170 dimensions, most low-signal.** xRFM must *learn* relevance via AGOP; trees ignore noise via split selection.
- **8.3% missingness.** All models receive median-imputed numerics in our pipeline, so we did not test native tree missing-value handling; a defensible hypothesis is that median imputation distorts similarity geometry more for kernels than for split-based models.

Choice of metrics. We report RMSE rather than MAE for regression because both Appliances Energy and Crop Mapping have heavy-tailed targets where large errors deserve disproportionate penalty; R^2 accompanies RMSE for cross-dataset comparability since target ranges differ by orders of magnitude (Appliances mean 97 Wh vs. Seoul Bike mean 705 rentals). For classification we report both Accuracy and AUC-ROC, but on IDA2016 (1.6% positive) AUC is the load-bearing number, since accuracy ≈ 0.984 is trivially achievable by always predicting "no failure", while AUC tests rank-ordering of failure probabilities, so the fact that all four models exceed 0.989 says they actually learn the rare class. Train/inference times in Tables 2–3 are wall-clock on shared cloud hardware with xRFM on GPU (CUDA) and the other three on CPU, so they reflect hardware asymmetry and should not be read as a method-fair complexity comparison.

AGOP vs. MI / permutation / PCA. Our random-feature test confirms §2.1's prediction that AGOP inherits model biases (tuned RFM demotes noise to rank 20–21/27; default-HP RFM only 12–13/27); MI and signed permutation sidestep this. Practical takeaway: **AGOP for interpreting a tuned model, MI or signed permutation as a model-free sanity check, PCA as a null baseline.**

With more time. Multi-seed replication (we used seed 42 only) would put error bars on close comparisons, especially HCC's 33-sample test. The HP search swept only `exponent` across two kernel slices (`l1`, `l2`); the full $K_{p,q}$ grid for $0 < q \leq p \leq 2$ is the natural follow-up. Smaller `max_leaf_size` sweeps would test the partitioned-tree regime our experiments only marginally entered.

5. Conclusion

xRFM is a credible complement to GBDTs in these experiments, especially where AGOP interpretability is useful. On our five UCI datasets it is competitive on smooth-target tasks (Appliances Energy, Crop Mapping), improves substantially with HP tuning on imbalanced high-dimensional classification (IDA2016), and trails on small-n (HCC) and rule-like targets (Seoul Bike) where tree inductive bias is better matched. However, the evidence is single-seed, only lightly exercises xRFM's tree partitioning, and is not sufficient to establish robust superiority over XGBoost or CatBoost.

References

1. Beaglehole, D., Holzmüller, D., Radhakrishnan, A., & Belkin, M. (2025). *xRFM: Accurate, scalable, and interpretable feature learning models for tabular data.* arXiv:2508.10053.
2. Radhakrishnan, A., Beaglehole, D., Pandit, P., & Belkin, M. (2024). Mechanism for feature learning in neural networks and backpropagation-free machine learning models. *Science*, **383**(6690).
3. Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In *KDD 2016*.
4. Breiman, L. (2001). Random Forests. *Machine Learning*, **45**(1), 5–32.
5. Prokhorenkova, L., Gusev, G., Vorobev, A., Dorogush, A. V., & Gulin, A. (2018). CatBoost: unbiased boosting with categorical features. In *NeurIPS 2018*.
6. Hollmann, N., Müller, S., Purucker, L., Krishnakumar, A., Körfer, M., Hoo, S. B., Schirrmeister, R. T., & Hutter, F. (2025). Accurate predictions on small data with a tabular foundation model. *Nature*, **637**.
7. Grinsztajn, L., Oyallon, E., & Varoquaux, G. (2022). Why do tree-based models still outperform deep learning on typical tabular data? In *NeurIPS 2022 Datasets & Benchmarks Track*.
8. Ye, H.-J., Liu, S.-Y., Rong, H.-R., Lu, P.-X., Du, J.-Y., Chao, W.-L., Wang, Z.-J., Sun, Q., & Jiang, Y. (2024). *A Closer Look at Deep Learning on Tabular Data (TALENT).* arXiv:2407.00956.
9. Constantine, P. G., Dow, E., & Wang, Q. (2014). Active subspace methods in theory and practice. *SIAM J. Sci. Comput.*, **36**(4).
10. Beaglehole, D., Súkeník, P., Mondelli, M., & Belkin, M. (2024). Average gradient outer product as a mechanism for deep neural collapse. In *NeurIPS 2024*.
11. Candanedo, L. M., Feldheim, V., & Deramaix, D. (2017). Data driven prediction models of energy use of appliances in a low-energy house. *Energy and Buildings*, **140**, 81–97.
12. Dua, D., & Graff, C. (2017). UCI Machine Learning Repository. University of California, Irvine.
13. Akiba, T., Sano, S., Yanase, T., Ohta, T., & Koyama, M. (2019). Optuna: A next-generation hyperparameter optimization framework. In *KDD 2019*.

Appendix A: Dataset Feature Descriptions

A.1 Seoul Bike Sharing (UCI 560)

Hourly rental counts in Seoul, Sept 2017 – Nov 2018 ($n = 8{,}760$, $d = 12$). Target: `Rented_Bike_Count` $\in [0, 3556]$, mean 704.6, std 645.0. **No missing values.**

- **Numeric (9):** `Hour` (0–23), `TemperatureC` (-17.8 to 39.4), `Humiditypct` (0–98), `Wind_speed_m/s` (0–7.4), `Visibility_10m` (27–2000), `Dew_point_temperatureC` (-30.6 to 27.2), `Solar_Radiation_MJ/m²` (0–3.52), `Rainfallmm` (0–35), `Snowfall_cm` (0–8.8).
- **Categorical (3):** `Seasons` (4 levels), `Holiday` (Yes/No), `Functioning_Day` (Yes/No).

The target follows clear rules: rentals are exactly zero whenever the bike system is off (`Functioning_Day=No`);

otherwise demand has two daily peaks (morning and evening commutes); and the time-of-day pattern depends strongly on the season. These are exactly the kind of axis-aligned, piecewise-constant patterns that decision trees handle naturally, which is a structural reason to expect GBDTs to do well here.

A.2 Appliances Energy Prediction (UCI 374)

10-minute energy readings from a Belgian house ($n = 19{,}735$, $d = 27$, all numeric). Target: `Appliances` (Wh) $\in [10, 1080]$, mean 97.7, std 102.5. **No missing values.**

- **Indoor sensors (18):** `T1`–`T9` (room temperatures, 14.9–29.9 °C), `RH_1`–`RH_9` (room humidities, 1–100%).
- **Outdoor weather (7):** `T_out`, `Press_mm_hg`, `RH_out`, `Windspeed`, `Visibility`, `Tdewpoint`, `lights`.
- **Random noise (2):** `rv1`, `rv2`, both drawn from $\mathcal{U}(0, 50)$ by Candanedo et al. (2017) as ground-truth-irrelevant features. Used as the "noise rejection" check in §3.2.

A.3 HCC Survival (UCI 423)

Hepatocellular carcinoma 1-year survival in a Portuguese cohort ($n = 165$, $d = 49$). Target: binary survival outcome, **class balance 102 lives : 63 dies (62%/38%)**. **826 missing values (10.2% of cells).** Features anonymised as `x0`–`x48`.

- **Categorical (27):** `x0`–`x22` are 22 binary clinical flags (cirrhosis, alcoholism, hepatitis, ascites, etc.); `x26` (5 levels: ECOG performance status), `x27`/`x28` (3 levels each: encephalopathy/ascites grades), `x43` (6 levels: Child-Pugh).
- **Numeric (22):** Lab values with widely varying scales: `x23` (age, 20–93), `x30` (likely AFP, range 1.2 to 1.8M), `x33`/`x34` (likely AST/ALT, max 13k and 459k respectively), plus haemoglobin, platelets, bilirubin, INR, creatinine, iron studies. Missingness up to 48% on rarely-ordered tests.

High $d/n \approx 0.30$, which is the report's primary hypothesis for xRFM's underperformance on this dataset.

A.4 IDA2016 Scania APS (UCI 414)

Heavy-truck Air Pressure System failure prediction ($n = 60{,}000$, $d = 170$, all numeric). Target: binary (failure/no), **class balance 59,000 : 1,000 (1.67% positive). 850,015 missing values (8.3% of cells).**

Features are anonymised sensor readings (`aa_000`, `ab_000`, …); the `ag_000`–`ag_009` series are histogram bins of one underlying counter. **Eight columns are >50% missing** (e.g., `br_000` 82%, `bq_000` 81%, `bp_000` 80%). All three properties stress kernel methods specifically. With 170 mostly-uninformative dimensions, xRFM has to *learn* which features matter, whereas trees ignore irrelevant features for free via split selection. With only 1.6% positive class, kernel ridge regression has no built-in mechanism

for asymmetric error costs, where GBDTs can be told to up-weight the minority. And with 8.3% of cells missing, although tree-based methods can in principle absorb missingness through their split logic, our pipeline imputes upstream (median for numerics, mode for categoricals) for all four models, so the kernel-vs-tree distinction here reduces to how well median imputation preserves geometry, not native missing-value handling. We expected this dataset to be the hardest of our five for xRFM, and it was.

A.5 Crop Mapping, Winnipeg (UCI 525)

Crop classification from fused satellite data over Manitoba, 2015–2016 ($n = 325{,}834$, stratified-subsampled to $n = 50{,}000$; $d = 174$, all numeric, no missing). Target: 7 crop classes, with **highly skewed prevalence** (largest class 26.1%; smallest 0.35%, only 175 samples).

Features come from two satellite systems collected at multiple dates: **RADARSAT-2 synthetic aperture radar** (4 polarisation channels : HH, HV, VH, VV; plus derived parameters) and **Landsat-8 optical bands** (visible, near-infrared, short-wave infrared). Raw values span roughly $[-28{,}400, 1{,}089{,}300]$; the radar measurements are in negative-dB units while optical radiances are large positive numbers, so standardisation is essential. Crucially for our hypothesis, all 174 features are continuous and physically correlated: the kind of smooth, structured feature space where a kernel's similarity-based assumption matches the data well.

Appendix B: Hyperparameter Tuning Details

B.1 Optuna setup. TPE sampler seeded with `random_state=42`. Validation objective: RMSE (minimise) for regression, AUC-ROC (maximise) for binary, accuracy (maximise) for multiclass. The configuration with the best validation objective is retrained on the train split and evaluated once on held-out test.

B.2 Search spaces and fixed knobs.

- **xRFM** (25 trials). Tunable: `bandwidth` $\in [1, 200]$ log, `exponent` $\in [0.7, 1.4]$, `kernel` $\in$ { `l1`, `l2` }, `reg` $\in [10^{-6}, 1]$ log, `diag` $\in \{T, F\}$. Fixed: `iters=5`, `M_batch_size=1000`, `bandwidth_mode=constant`, `max_leaf_size=min(60000, n)`.
- **XGBoost** (40 trials). Tunable: `max_depth` $\in [3, 12]$, `learning_rate` $\in [10^{-3}, 0.3]$ log, `subsample` $\in [0.5, 1]$, `colsample_bytree` $\in [0.5, 1]$, `reg_alpha` $\in [10^{-8}, 10]$ log, `reg_lambda` $\in [10^{-8}, 10]$ log, `min_child_weight` $\in [1, 10]$. Fixed: `n_estimators=2000` with `early_stopping_rounds=50`.
- **CatBoost** (30 trials). Tunable: `depth` $\in [4, 10]$, `learning_rate` $\in [10^{-3}, 0.3]$ log, `l2_leaf_reg` $\in [10^{-2}, 10]$ log, `bagging_temperature` $\in [0, 1]$, `border_count` $\in \{32, 64, 128, 254\}$. Fixed: `iterations=2000` with `early_stopping_rounds=50`.
- **Random Forest** (20 trials). Tunable: `n_estimators` $\in [200, 1000]$, `max_depth` $\in [5, 30]$, `min_samples_split` $\in [2, 20]$, `min_samples_leaf` $\in [1, 10]$, `max_features` $\in$ { `sqrt`, `log2`, 0.5 }.

B.3 Selected hyperparameters per dataset. Best configurations chosen by Optuna's TPE search (Optuna [Akiba et al., 2019] is a hyperparameter tuning library; its TPE sampler picks each next trial near configurations that scored well on previous trials, which makes it more sample-efficient than random or grid search). Configurations extracted from `results/main_*_seed42.json`. RF abbreviations: `mss`=`min_samples_split`, `msl`=`min_samples_leaf`.

xRFM:

- Appliances Energy: `bw`=12.72, `exp`=0.79, `kernel`=l1, `reg`=0.061, `diag`=T
- Seoul Bike: `bw`=2.09, `exp`=0.90, `kernel`=l1, `reg`=0.051, `diag`=F
- HCC Survival: `bw`=23.08, `exp`=0.73, `kernel`=l2, `reg`=2.5×10^{-6}, `diag`=F
- IDA2016: `bw`=199.3, `exp`=1.27, `kernel`=l2, `reg`=0.95, `diag`=T
- Crop Mapping: `bw`=4.11, `exp`=1.18, `kernel`=l2, `reg`=5.4×10^{-6}, `diag`=T

XGBoost:

- Appliances Energy: `depth`=12, `lr`=0.042, `sub`=0.98, `col`=0.51, =0.007, =1.48, `mcw`=2
- Seoul Bike: `depth`=7, `lr`=0.007, `sub`=0.62, `col`=0.99, $=3\times10^{-6}$, $=4\times10^{-4}$, `mcw`=2
- HCC Survival: `depth`=5, `lr`=0.253, `sub`=0.83, `col`=0.78, =0.026, =0.14, `mcw`=2
- IDA2016: `depth`=8, `lr`=0.004, `sub`=0.87, `col`=0.75, =0.41, =9.65, `mcw`=10
- Crop Mapping: `depth`=3, `lr`=0.160, `sub`=0.69, `col`=0.72, $=6\times10^{-6}$, $=2\times10^{-5}$, `mcw`=5

CatBoost:

- Appliances Energy: `depth`=9, `lr`=0.051, `l2`=0.066, `bag`=0.35, `border`=254
- Seoul Bike: `depth`=8, `lr`=0.031, `l2`=0.115, `bag`=0.36, `border`=64
- HCC Survival: `depth`=8, `lr`=0.057, `l2`=0.012, `bag`=0.97, `border`=32
- IDA2016: `depth`=9, `lr`=0.137, `l2`=4.37, `bag`=0.85, `border`=64
- Crop Mapping: timed out (see §3.1).

Random Forest:

- Appliances Energy: `n`=652, `depth`=24, `mss`=10, `msl`=1, `max_feat`=sqrt
- Seoul Bike: `n`=422, `depth`=25, `mss`=12, `msl`=3, `max_feat`=0.5
- HCC Survival: `n`=796, `depth`=24, `mss`=17, `msl`=1, `max_feat`=log2
- IDA2016: `n`=599, `depth`=24, `mss`=5, `msl`=7, `max_feat`=log2
- Crop Mapping: `n`=652, `depth`=24, `mss`=10, `msl`=1, `max_feat`=sqrt